

Acte

Project Brief & Functional Specification — for Darwin (Development)
v2 — based on the working HTML prototype (“acte-dashboard-v16.html”)

Acte is a time-tracking & billing product for law firms (“cabinets d'avocats”). An AI-powered background tracker passively captures what each lawyer does on their computer (Word, Outlook, browser) and automatically converts it into billable time entries linked to the right case file, with a confidence score. Lawyers validate/correct entries; the firm's admin manages the team and subscription. A working front-end prototype already exists (Yann's mockup) and is summarized below so it can be used as the functional reference for the real build.

1. Product Structure

The product has two sides:

- **Lawyer / user interface** — the daily tool each lawyer/staff member at the firm uses (Home, Time Journal, Case Files, Billing, Profile & Impact, Cloud & Sync, Settings).
- **Admin Console (“Cabinet”)** — for Yann (or any partner/admin) to manage the firm's users and subscription. Reached from the profile menu.

In the prototype, both live in the same app and the admin simply switches view (no separate login) — confirm with Darwin whether the real product should also gate the Admin Console by role/permission.

2. Lawyer / User Interface

2.1 Home

- Time captured today, split into *validated* vs *pending* minutes.
- Secured revenue this month (“CA sécurisé”) and average hourly rate.
- “ACTE ROI” widget: minutes recovered today vs. manual time entry.
- AI panel (“Le Cerveau d'ACTE”): short natural-language daily summary / insight.
- List of pending tasks awaiting validation, with a “Validate all” (“Tout intégrer”) batch action.
- Weekly billed-time summary chart.

2.2 Time Journal

The core screen: a list of tasks automatically captured by the AI tracker during the day. Each task row shows:

- Title (auto-generated, e.g. “Drafting submissions — Delcourt v. Mutuelle Azur”).
- Source badge: **Word** / **Outlook** / **Browser**.
- Time range and duration (e.g. 09:15–10:32 · 1h17).
- Case file (“dossier”) it's attributed to — editable via a dropdown of the firm's case files.
- Confidence score badge (see 5.2 below) — color-coded, with an “AI reasoning” (“why”) explanation of why the task was matched to that case.
- A validate button, and the option to reassign the case, which is logged so the AI “learns” from the correction.

Example captured task (from the prototype's mock data):

Source	Title	Duration	Confidence	Case file
Word	Drafting submissions — Delcourt v. Mutuelle Azur	1h17	97%	Delcourt v. Mutuelle Azur
Browser	Case-law research — disproportionate guarantee	23 min	84%	Vasseur — Guarantee
Outlook	Settlement email — Mr. Bône	15 min	92%	Bône v. SCI Alma
Word	Reviewing commercial lease — SARL Toison d'Or	27 min	76% — needs review	SARL Toison d'Or — Lease

2.3 Case Files (“Dossiers”)

- Grid of active case files: name, client, hours used vs. an hour budget, status (*in progress / ready to bill / archived*), last activity timestamp.
- Create a new case file (name, client, hour budget) from a modal; it immediately becomes available in the Journal's case-assignment dropdown.
- Edit a case's hour budget; the AI is expected to warn the user as usage approaches the budget limit.
- Archive / restore a case file (archiving suspends time capture for that case).

2.4 Billing

- View of validated time that is ready to be invoiced, and a weekly billed-time summary.

2.5 Profile & Impact

- Personal AI-generated monthly report, e.g. hours of “invisible”/volatile time captured in the background and the corresponding revenue secured.
- This is also where the user's tracker activation key lives (see section 4.2).

2.6 Cloud & Synchronization

- List of the user's linked devices (e.g. office PC, personal MacBook), each with the Companion tracker version, online/offline sync status, and last-sync time.
- Force a manual sync; unlink a device; export synced data.

3. Admin Console (“Cabinet”) — for Yann

Two tabs: **Team** (Équipe du cabinet) and **Firm Subscription** (Abonnement Cabinet).

3.1 Team tab

Per-member fields:

- Name, role (Partner / Associate / Paralegal / Trainee lawyer, etc.).
- Minutes tracked and validation rate (%) over the period.
- Status: *active / in court / suspended / invitation pending*.
- Partner flag and hourly billing rate (role sets a default rate, editable per person).

Admin actions on a member:

- Edit profile / hourly rate.

- Send a reminder to validate their pending time entries.
- Suspend / reactivate the account.
- Invite a new member (pick role → default rate is pre-filled) — sends an invitation; admin can resend or cancel a pending invite.

Cabinet-level AI insight example from the prototype: flags the team member with the lowest validation rate, plus aggregate stats (total hours captured, global validation rate, number of active members).

3.2 Firm Subscription tab

- Seat-based pricing: price = number of seats × €50/month, no commitment, shown live as seats change.
- Billing handled via Stripe in the intended real product (the prototype only simulates this).
- Invoice history with downloadable invoices.

3.3 Security messaging shown to the firm

The prototype already states a specific privacy/security architecture to the firm's admin — this is effectively a product requirement, not just copy:

- Actual captured content (documents, emails, browsing) is processed **locally** on each member's device.
- Only time **metadata** (duration, case file, source app) is sent to the console.
- That metadata is encrypted (**AES-256**) in transit/at rest.
- Professional confidentiality (“secret professionnel”) stays under the firm's exclusive control.
- ISO 27001 certification “in progress”, annual security audits — positioning only for now, confirm real compliance targets with Yann before it ships as a claim.

4. Activity Tracker (“ACTE Tracker” / “Compagnon”)

This is the background component that observes each user's computer and feeds the Time Journal. In the prototype it's presented as a downloadable desktop “Companion” app:

4.1 Distribution (as mocked)

- Windows 10/11 — ACTE_Tracker .exe (~38 MB).
- macOS 13+ — ACTE_Tracker .dmg (~42 MB).
- Signed installer, OS auto-detected on the download page, versioned (e.g. v1.4.2).
- A user can link several devices (e.g. office PC + personal Mac), each tracked separately in Cloud & Sync.

4.2 Linking / activation

Each user gets a unique activation/API key (format seen in the prototype: ACTE-VC-7F42-K9D1-3MB8-QX5E), shown on their Profile page. The key is pasted into the desktop Companion app to link that specific workstation to their account, and can be revoked at any time from the web app.

4.3 What it tracks (sources to start with)

- **Word** — which document is active and for how long (matched to a case via filename/content cues).
- **Outlook** — emails sent/received, with the correspondent matched against known case contacts.
- **Browser** — research sessions on specific reference legal-research sites. The prototype uses **Légifrance** and **Dalloz** as placeholder examples.

Reference site list: the definitive list of sites lawyers should be tracked on is **still pending from Yann**. Légifrance/Dalloz in the mockup are illustrative only, not confirmed scope.

4.4 Open engineering question for Darwin

The prototype only mocks the *concept* (a downloadable Companion + activation key). It does not define how the tracker actually runs reliably in the background on each user's machine: background service/daemon vs. browser extension vs. OS-level hooks, required permissions, auto-start on login, update mechanism, behavior across Windows/macOS, and how it stays lightweight and unobtrusive. This needs a real technical design from Darwin.

4.5 Privacy architecture (already defined, treat as a constraint)

Per section 3.3: raw content must stay local to the device; only time metadata (duration, case, source) should ever reach the server, encrypted. This should shape the tracker's design from the start rather than being retrofitted.

5. AI Case-Matching & Confidence Scoring

For every captured task, the AI (1) assigns it to a case file, and (2) attaches a confidence score, plus a short human-readable explanation of the reasoning (e.g. “email sent to contact@sci-alma.fr, a known contact on the Bône v. SCI Alma case”).

5.1 Correction feedback loop

If the user reassigns a task to a different case, the system should log that correction (the prototype shows a message like “corrected — I'll take it into account”). Whether this actually retrains/adjusts future matching, or is just an audit log for now, needs to be defined with Darwin.

5.2 Confidence score thresholds (UI behavior already defined)

Confidence	Display	Meaning
≥ 90%	Gold badge	High confidence, no action needed
80–89%	Neutral badge	Acceptable, no action needed
< 80%	Amber badge — “needs review”	Should be flagged / trigger an alert

This matches the requirement of alerting when a task is below 80% reliability — in the prototype it's a visual badge; it should also drive an actual notification (see section 6). The underlying scoring algorithm itself (how the % is computed) is not defined anywhere — it's mock data in the prototype and needs to be designed.

6. Alerts & Notifications

Required alert types:

- **Low-confidence task** — a captured task scores below 80% and needs the user's review (already visualized in the Journal; needs to become a real notification, not just a badge).
- **Case budget approaching limit** — a case file's tracked hours are nearing its hour budget.
- **Team validation lagging** (admin-facing) — a team member has an unusually low validation rate or a backlog of unvalidated time.
- **General app/account health** — sync issues (device offline), pending invitations, subscription/billing issues, or anything else relevant to managing the app.

Delivery channel (in-app only, email, or push from the desktop Companion) is not yet defined and should be discussed with Darwin.

7. Feedback Page (like Epidom)

The requirement is a dedicated feedback page in the Admin Console, similar to the one built for Epidom. The current prototype does **not** have this — it only has an internal AI activity feed/log (a running log of validations and corrections, for context/audit, not user-submitted feedback). This page still needs to be designed and built from scratch, reusing the Epidom pattern.

8. Open Points / Pending

- Reference list of legal-research websites to track — pending from Yann (Légifrance/Dalloz are placeholders in the mockup only).
- Technical approach for running the tracker reliably in the background on each user's machine — to be defined with Darwin (see 4.4).
- Definition of the confidence-scoring algorithm behind the AI's case-matching (currently mock data only).
- Alert delivery channel(s) — in-app, email, desktop notification.
- Whether the Admin Console needs its own role/permission gate, separate from the regular user view.
- The dedicated Feedback page (Epidom-style) still needs to be designed — not present in the current prototype.
- Real compliance status behind the ISO 27001 / security-audit claims shown to firms.

Prepared by Evan — Prionation, based on Yann's working prototype ("acte-dashboard-v16.html", 8 July 2026 build). This brief reflects what the prototype currently does; it will evolve as open questions are resolved with Yann and Darwin.